

# RatuDB

write by: lei.fu

?????

????? Opensearch 3.0 ?????????????? Opensearch????? Opensearch ??????Cassandra?

Cassandra????????Elasticsearch????Cassandra????????????????????  
????????????Cassandra????????

Opensearch

?? ???? ????RatuDB????????.

## Getting Started

### Installation

- [??] ?? ?? RatuDB ???.
- ?? `bin/opensearch` ? Linux ?? macOS. ?? `bin\opensearch.bat` ? Windows.
- ?? `curl -X GET http://localhost:9200/` ?? Elasticsearch ???.
- ????????? `cql` ???, ?????????, ?????????  
`./bin/cqlsh` `cql`?????`ip`?????????????`127.0.0.1`?????Cassandra?????
- ??????`./bin/nodetool` ??????????????RatuDB????? 9042????????????????????????????????

### 1.????

????????????????Cassandra?

```
$ ./bin/cqlsh
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.1.3-SNAPSHOT | CQL spec 3.4.6 | Native protocol v5]
Use HELP for help.
cqlsh>
```

????????keyspaces:

```
CREATE KEYSPACE lei
WITH REPLICATION = {'class': 'SimpleStrategy', 'replication_factor': 1};
```

`replication_factor` ??????????????Cassandra????????????????????

□ □ □ □ □ □ □ :

```
CREATE TABLE lei.tweets (
  id INT PRIMARY KEY,
  user TEXT,
  body TEXT,
  time TIMESTAMP,
  latitude FLOAT,
  longitude FLOAT
);
```

## 2.□□□□

**Cassandra**

```
CREATE CUSTOM INDEX tweets_index ON lei.tweets ()
USING 'org.apache.ratu.second.ElasticSecondaryIndex'
WITH OPTIONS = {
  'refresh_seconds': '1',
  'async_write': 'true',
  'schema': '{
    fields: {
      id: {type: "integer"},
      user: {type: "text"},
      body: {type: "text", analyzer: "english"},
      time: {type: "date", pattern: "yyyy-MM-dd"},
      latitude: {type: "float"},
      longitude: {type: "float"}
    }
  }'
```

□ □

[illegible]

```
CREATE CUSTOM INDEX tweets_index ON lei.tweets () -- tweets_index [] lei.tweets
-- []
```

```
USING 'org.apache.ratu.second.ElasticSecondaryIndex' ████████████████████████
```

```
WITH OPTIONS (refresh_seconds = 1) Opensearch
ES
ES
ES
`-1` Opensearch api
```

POST / refresh

schema 的 fields 是 OpenSearch 的

body: {type: "text", analyzer: "english"}, 的 type 是 analyzer

time: {type: "date", pattern: "yyyy-MM-dd"}, 的 pattern

的 text

es



### 3.

1G

```
[OVERALL], RunTime(ms), 4167
[OVERALL], Throughput(ops/sec), 2399.808015358771
[TOTAL_GCS_G1_Young_Generation], Count, 4
[TOTAL_GC_TIME_G1_Young_Generation], Time(ms), 16
[TOTAL_GC_TIME_%_G1_Young_Generation], Time(%), 0.3839692824574034
[TOTAL_GCS_G1_Old_Generation], Count, 0
[TOTAL_GC_TIME_G1_Old_Generation], Time(ms), 0
[TOTAL_GC_TIME_%_G1_Old_Generation], Time(%), 0.0
[TOTAL_GCs], Count, 4
[TOTAL_GC_TIME], Time(ms), 16
[TOTAL_GC_TIME_%], Time(%), 0.3839692824574034
[CLEANUP], Operations, 10
[CLEANUP], AverageLatency(us), 222516.6
[CLEANUP], MinLatency(us), 0
[CLEANUP], MaxLatency(us), 2226175
[CLEANUP], 95thPercentileLatency(us), 2226175
[CLEANUP], 99thPercentileLatency(us), 2226175
[INSERT], Operations, 10000
[INSERT], AverageLatency(us), 1289.455
[INSERT], MinLatency(us), 401
[INSERT], MaxLatency(us), 47551
[INSERT], 95thPercentileLatency(us), 3957
[INSERT], 99thPercentileLatency(us), 5307
[INSERT], Return=OK, 10000
```

```
[OVERALL], RunTime(ms), 4375
[OVERALL], Throughput(ops/sec), 2285.714285714286
[TOTAL_GCS_G1_Young_Generation], Count, 4
[TOTAL_GC_TIME_G1_Young_Generation], Time(ms), 15
[TOTAL_GC_TIME_%_G1_Young_Generation], Time(%), 0.34285714285714286
[TOTAL_GCS_G1_Old_Generation], Count, 0
[TOTAL_GC_TIME_G1_Old_Generation], Time(ms), 0
[TOTAL_GC_TIME_%_G1_Old_Generation], Time(%), 0.0
[TOTAL_GCs], Count, 4
[TOTAL_GC_TIME], Time(ms), 15
[TOTAL_GC_TIME_%], Time(%), 0.34285714285714286
[CLEANUP], Operations, 10
[CLEANUP], AverageLatency(us), 223336.7
[CLEANUP], MinLatency(us), 1
[CLEANUP], MaxLatency(us), 2234367
[CLEANUP], 95thPercentileLatency(us), 2234367
[CLEANUP], 99thPercentileLatency(us), 2234367
[INSERT], Operations, 10000
[INSERT], AverageLatency(us), 1575.29
[INSERT], MinLatency(us), 541
[INSERT], MaxLatency(us), 56607
[INSERT], 95thPercentileLatency(us), 3393
[INSERT], 99thPercentileLatency(us), 7135
[INSERT], Return=OK, 10000
```

□ □ □ □ □ □ □ □ □ □ □

- `ES` `Queue` `Queue` `[parent]` `Data` `too large`` `ES`
- `async_write`` `false`` `true``

□ □ □ □ □ □ □

```
CREATE CUSTOM INDEX usertable_index ON ycsb.usertable ()
USING 'org.apache.ratu.second.ElasticSecondaryIndex'
WITH OPTIONS = {
    'refresh_seconds': '120',
    'async_write': 'true',
    'schema': '{
        fields: {
            field0: {type: "text"},
            field1: {type: "text"},
            field2: {type: "text"},
            field3: {type: "text"},
            field4: {type: "text"},
            field5: {type: "text"},
            field6: {type: "text"},

```

```

        field7: {type: "text"},
        field8: {type: "text"},
        field9: {type: "text"}
    }
}';
};

```

es es es queue api:

```
GET _cat/thread_pool?v
```

queue es

refresh\_seconds` -

1`Opensearch`[parent] Data too large`  
OpensearchJVM`jvm-  
options`Opensearch

## 4.

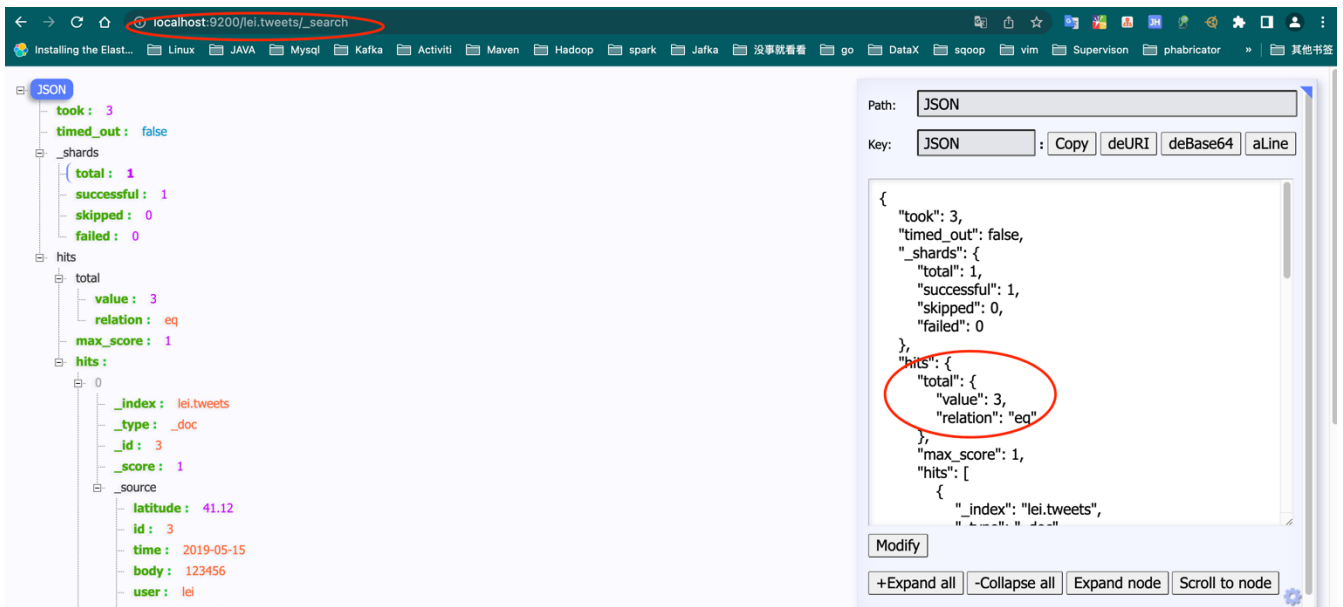
,

```
INSERT INTO lei.tweets (id, user, body, time,latitude,longitude) VALUES (1, 'fu',
'abc', '2015-05-15',41.12,-71.34);
```

```
INSERT INTO lei.tweets (id, user, body, time,latitude,longitude) VALUES (2, 'fu',
'123456', '2019-05-15',41.12,-71.34);
```

```
INSERT INTO lei.tweets (id, user, body, time,latitude,longitude) VALUES (3, 'lei',
'123456', '2019-05-15',41.12,-71.34);
```

:



□□□□□□

## 5.□□

□□□□□□□□□□Cassandra□□□□□□□□□□□□□□□□□□□□□□□□□□□□

```
<!--range-->
SELECT * FROM lei.tweets WHERE expr(tweets_index, '{
  query: {type: "range", field: "time", gte: "2014-04-25", lte: "2015-05-21"}
}');
```

```
query:□□□□□□□□
type:□□□□□DSL□□□□□
field: □□□□□□□□□□
gte: □□□□
lte: □□□□
```

```
<!--match-->
SELECT * FROM lei.tweets WHERE expr(tweets_index, '{
  query: {type: "match", field: "user", query: "lei"}
}');
```

```
<!--match, value-->
SELECT * FROM lei.tweets WHERE expr(tweets_index, '{
  query: {type: "match", field: "user", value: "lei"}
}');
```

```
<!--match_phrase-->
```

```
SELECT * FROM lei.tweets WHERE expr(tweets_index, '{
  query: {type: "match_phrase", field: "user", query: "lei"}
}');
```

```
<!--match_phrase, value-->
SELECT * FROM lei.tweets WHERE expr(tweets_index, '{
  query: {type: "match_phrase", field: "user", value: "lei"}
}');
```

```
<!--term -->
SELECT * FROM lei.tweets WHERE expr(tweets_index, '{
  query: {type: "term", field: "user", value: "lei"}
}');
```

refresh: true ES ES

```
<!--range-->
SELECT * FROM lei.tweets WHERE expr(tweets_index, '{
  query: {type: "range", field: "time", gte: "2014-04-25", lte: "2015-05-21"},
  refresh: true
}') limit 100;
```

## Cassandra ↔ Opensearch

CQL	Java	ES	
ascii	String	text	ascii
bigint	long	long	64
blob	ByteBuffer/byte[]	text	ES
boolean	Boolean	boolean	
decimal	BigDecimal	float	
double	double	double	64
float	float	float	32
inet	String	ip	ipv4ipv6ip(ipv6)
int	int	integer	32
text	String	text	utf-8

CQL 名	Java 名	ES 名	備 考
timestamp	Date	date	Opensearch yyyy-MM-dd yyyy-MM- ddTHH:MM:SSZ ,
uuid	UUID	text	UUID
timeuuid	UUID	text	UUID
varchar	string	text	text
varint	BigInteger	text	
duration	String	text	
smallint	Integer	integer	16
tinyint	Integer	integer	8
list<T>	String	text	ES array
time	long	long	hh:mm:ss ES64
set<T>	Set<T>	text	ES array
map<T,T>	Map<T,T>	nested	

## 環境構築

RatuDB を [Gradle](#) でビルド。

Gradle を [8.0.2](#) から Gradle を

[distributions/archives](#) から。

## Fork

RatuDB を Cassandra でビルド。

```
git clone ssh://git@gitlab.ratu.ltd:30022/operation-ratudb/ratudb-opensearch.git
```

Cassandra を [submodule](#) で管理。

```
git clone ssh://git@gitlab.ratu.ltd:30022/operation-ratudb/ratudb-opensearch.git
--recursive
```



```
./gradlew run
```

#####JDK11, #####RatuDB#####IDEA#####

```
.gradlew (java)
[2023-08-02T10:56:43,035][INFO ][o.a.c.d.c.CompactionTask ] [runTask-0] Compacted (35b22920-30e0-11ee-a5a4-350b91546892)
6 sstables to [/Users/lei.fu/work/ratudb-es/build/testclusters/runTask-0/data/data/system/local-7ad54392bccdd35a684174e047
860b377/nb-7-big,] to level=0. 0.819KiB to 0.686KiB (~83% of original) in 354ms. Read Throughput = 2.312KiB/s, Write Th
roughput = 1.934KiB/s, Row Throughput = ~2/s. 6 total partitions merged to 1. Partition merge counts were {6:1, }. Time
spent writing keys = 43ms
[2023-08-02T10:56:43,036][INFO ][o.a.c.i.s.SSTable ] [runTask-0] Deleting sstable: /Users/lei.fu/work/ratudb-es/bu
ild/testclusters/runTask-0/data/data/system/local-7ad54392bccdd35a684174e047860b377/nb-1-big
[2023-08-02T10:56:43,039][INFO ][o.a.c.i.s.SSTable ] [runTask-0] Deleting sstable: /Users/lei.fu/work/ratudb-es/bu
ild/testclusters/runTask-0/data/data/system/local-7ad54392bccdd35a684174e047860b377/nb-2-big
[2023-08-02T10:56:43,041][INFO ][o.a.c.i.s.SSTable ] [runTask-0] Deleting sstable: /Users/lei.fu/work/ratudb-es/bu
ild/testclusters/runTask-0/data/data/system/local-7ad54392bccdd35a684174e047860b377/nb-3-big
[2023-08-02T10:56:43,043][INFO ][o.a.c.i.s.SSTable ] [runTask-0] Deleting sstable: /Users/lei.fu/work/ratudb-es/bu
ild/testclusters/runTask-0/data/data/system/local-7ad54392bccdd35a684174e047860b377/nb-6-big
[2023-08-02T10:56:43,045][INFO ][o.a.c.i.s.SSTable ] [runTask-0] Deleting sstable: /Users/lei.fu/work/ratudb-es/bu
ild/testclusters/runTask-0/data/data/system/local-7ad54392bccdd35a684174e047860b377/nb-5-big
[2023-08-02T10:56:43,047][INFO ][o.a.c.i.s.SSTable ] [runTask-0] Deleting sstable: /Users/lei.fu/work/ratudb-es/bu
ild/testclusters/runTask-0/data/data/system/local-7ad54392bccdd35a684174e047860b377/nb-4-big
[2023-08-02T10:56:52,770][WARN ][o.a.c.u.FBUtilities ] [runTask-0] Trigger directory doesn't exist, please create it
and try again.
[2023-08-02T10:56:52,774][INFO ][o.a.c.a.CassandraRoleManager] [runTask-0] Created default superuser role 'cassandra'
<=====--> 99% EXECUTING [2m 17s]
> :modules:repository-url:bundlePlugin
> :server:compileJava11Java
> :modules:percolator:compileJava
> :modules:repository-url:compileJava
> :modules:transport-netty4:compileJava
> :x-pack:plugin:ilm:compileJava
> :modules:lang-expression:compileJava > Unpacking build cache entry
> :modules:aggs-matrix-stats:compileJava
```

#####Cassandra#####

## 1. JDK

#####JDK

```
export RUNTIME_JAVA_HOME="JDK14"
export CASSANDRA_USE_JDK11="jdk11"
export JAVA_HOME="JDK11"
```

RUNTIME\_JAVA\_HOME 是 JDK CASSANDRA\_USE\_JDK11 Cassandra JDK

## 2.

#####Cassandra ant #####RatuDB#####Cassandra api

#####Cassandra api#####Cassandra#####

#####Cassandra Api

```
./gradlew cassandra-mvn-install
```

#####Cassandra api

□□□□□□

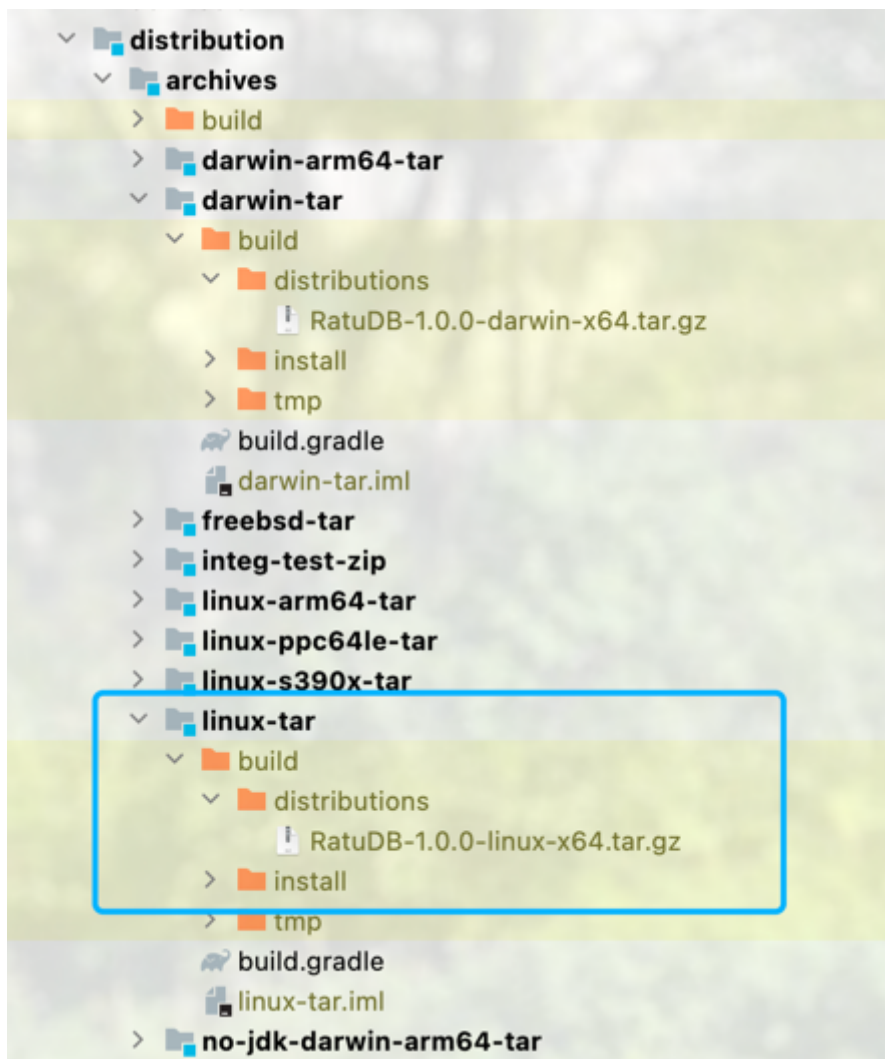
```
./gradlew assemble
```

Docker
 Docker

## 1. linux

```
./gradlew :distribution:archives:linux-tar:assemble
```

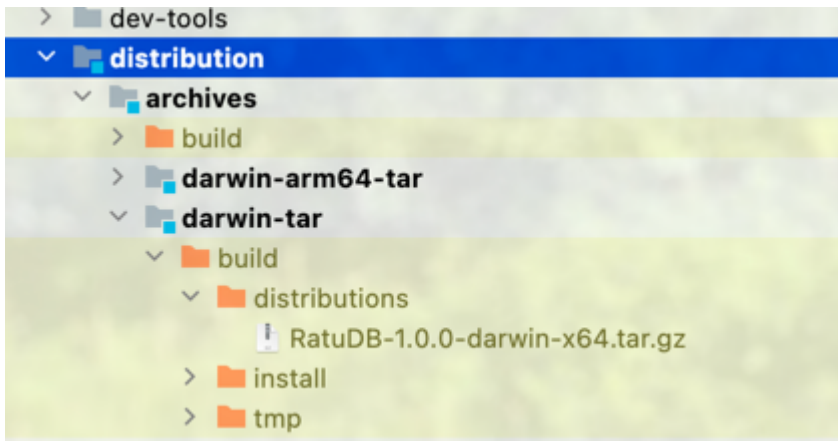
linux 00000000



## 2. mac

```
./gradlew :distribution:archives:darwin-tar:assemble
```

mac 00000000



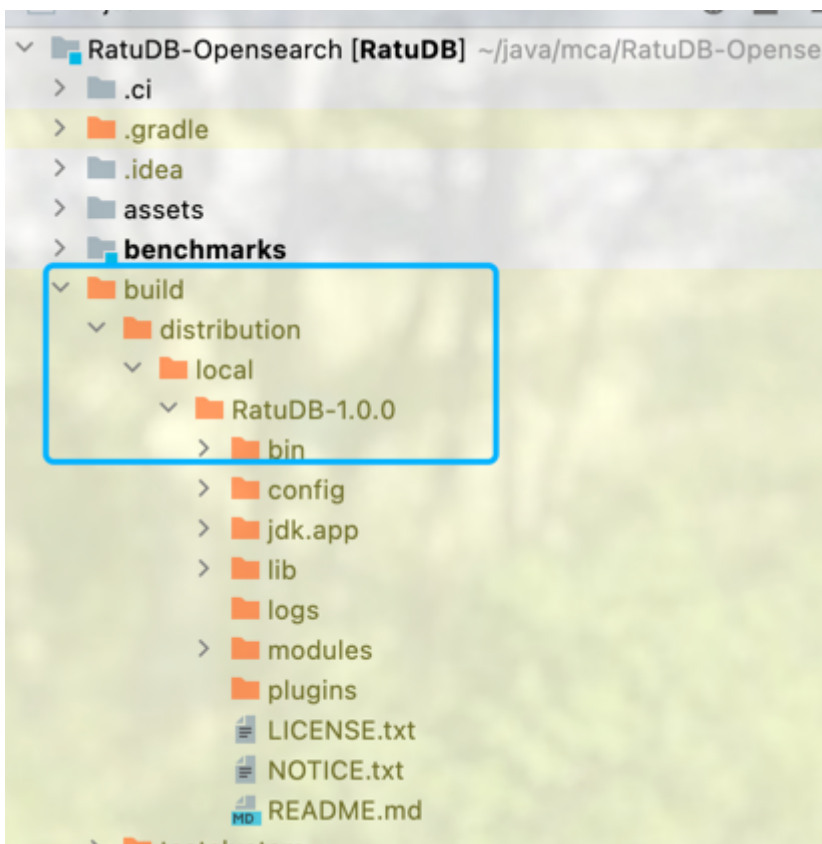
### 3. windows

```
./gradlew :distribution:archives:windows-zip:assemble
```

~~~~~

```
./gradlew localDistro
```

~~~~~mac~~~~~mac~~~~~



## IntelliJ IDEA

RatuDB JDK11, JDK11

- **File > Open**
- `build.gradle`
- **Open as Project**

□ □ □ □

111

## 1. linux

```
192.168.184.31
192.168.184.32
192.168.184.33
```

```
JDK 11 linux jdk jdk adoptopenjdk
15 Cassandra jdk linux
```

```
cat /etc/profile
```

```
export JAVA_HOME=/home/elastic/jdk-11
export PATH=$JAVA_HOME/bin:$PATH
export CLASSPATH=.:$JAVA_HOME/lib/dt.jar:$JAVA_HOME/lib/tools.jar
```

## 2. liunx

```
./gradlew :distribution:archives:linux-tar:assemble
```

### 3. RatuDB

#### 4.000000; RatuDB config cassandra.yaml

□ □ □ □ □ □ □ □

```
cluster_name: 'Ratu Cluster'
```

Cassandra Cassandra

```
rpc_address: 192.168.184.31
```

es`network.host`

```
listen_address: 192.168.184.31
```

□ □ □ □ □

```
- seeds: "192.168.184.31:7000,192.168.184.32:7000,192.168.184.33:7000"
```

```

#####ES#   discovery.seed_hosts   #   cluster.initial_master_nodes   #####http.port   #
transport.port #####9200 # 9300 #

```

## 5. cassandra-rackdc.properties

```
dc=dc1
rack=rack1
```

dc

rack

ES

History Settings Help

200 - OK 14 ms

1 GET \_cat/nodeattrs?v

|    | node           | host           | ip             | attr            | value     |
|----|----------------|----------------|----------------|-----------------|-----------|
| 2  | 192.168.184.33 | 192.168.184.33 | 192.168.184.33 | rack_id         | dc1-rack1 |
| 3  | 192.168.184.33 | 192.168.184.33 | 192.168.184.33 | xpack.installed | true      |
| 4  | 192.168.184.33 | 192.168.184.33 | 192.168.184.33 | transform.node  | true      |
| 5  | 192.168.184.32 | 192.168.184.32 | 192.168.184.32 | rack_id         | dc1-rack1 |
| 6  | 192.168.184.32 | 192.168.184.32 | 192.168.184.32 | xpack.installed | true      |
| 7  | 192.168.184.32 | 192.168.184.32 | 192.168.184.32 | transform.node  | true      |
| 8  | 192.168.184.31 | 192.168.184.31 | 192.168.184.31 | rack_id         | dc1-rack1 |
| 9  | 192.168.184.31 | 192.168.184.31 | 192.168.184.31 | xpack.installed | true      |
| 10 | 192.168.184.31 | 192.168.184.31 | 192.168.184.31 | transform.node  | true      |
| 11 |                |                |                |                 |           |

□ □ □ □ □ □ □ □ □ □ □ □ □ □ □ □ □ □ □ □

## 6. cassandra.yaml の data\_file\_directories を cassandra/es の data に変更

[[parent] Data too large]

Opensearch [parent] JVM 95%

JVM

□ □ □ □ □ □ □

```
GET _cat/nodes?v=true&h=name,node*,heap*
```

API

```
GET _nodes/stats/breaker
```

□ □ □ □ □ □ □ □ □ □ □ □ □ □ □ □

## Elasticsearch

1 Elasticsearch queue FIFO 10 100

Cassandra Cassandra commitLog Cassandra

2 translog es translog FileSystemCache segment translog

```
PUT /_settings?preserve_existing
{
  "index.translog.durability" : "async",
  "index.translog.flush_threshold_size" : "1024mb",
  "index.translog.sync_interval" : "60s"
}
```

20%